

EXPERIMENT:11

AIM:Create a R program to implement CART, K-Nearest Neighbours, SVM, LDA, Random Forest on PimaIndianDiabetes dataset and compare and visualize the results

DESCRIPTION:The “PimaIndianDiabetes” Dataset, originally from the National Institute of Diabetes and Digestive and Kidney Diseases, contains information of 768 women from a population near Phoenix, Arizona, USA. The outcome tested was Diabetes, 258 tested positive and 500 tested negative. Therefore, there is one target (dependent) variable and the 8 attributes (TYNECKI, 2018): pregnancies, OGTT(Oral Glucose Tolerance Test), blood pressure, skin thickness, insulin, BMI(Body Mass Index), age, pedigree diabetes function.

CODE:**# load libraries**

```
library(mlbench)
```

```
library(caret)
```

load the dataset

```
data(PimaIndiansDiabetes)
```

prepare training scheme

```
control <- trainControl(method="repeatedcv", number=10, repeats=3)
```

CART

```
set.seed(7)
```

```
fit.cart<- train(diabetes~., data=PimaIndiansDiabetes, method="rpart", trControl=control)
```

LDA

```
set.seed(7)
```

```
fit.lda<- train(diabetes~., data=PimaIndiansDiabetes, method="lda", trControl=control)
```

SVM

```
set.seed(7)
```

```
fit.svm<- train(diabetes~., data=PimaIndiansDiabetes, method="svmRadial", trControl=control)
```

kNN

```
set.seed(7)
```

```
fit.knn<- train(diabetes~., data=PimaIndiansDiabetes, method="knn", trControl=control)
```

Random Forest

```
set.seed(7)
```

```
fit.rf<- train(diabetes~., data=PimaIndiansDiabetes, method="rf", trControl=control)
```

collect resamples

```
results <- resamples(list(CART=fit.cart, LDA=fit.lda, SVM=fit.svm, KNN=fit.knn, RF=fit.rf))
```

summarize differences between modes

```
summary(results)
```

OUTPUT:

```
Call:
summary.resamples(object = results)

Models: CART, LDA, SVM, KNN, RF
Number of resamples: 30

Accuracy
```

	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.	NA's
CART	0.6753247	0.7272727	0.7532468	0.7469697	0.7662338	0.7922078	0
LDA	0.7142857	0.7508117	0.7662338	0.7791069	0.8000256	0.9078947	0
SVM	0.7236842	0.7508117	0.7631579	0.7712919	0.7915243	0.8947368	0
KNN	0.6753247	0.7036056	0.7272727	0.7369503	0.7662338	0.8311688	0
RF	0.6842105	0.7305195	0.7597403	0.7638528	0.8019481	0.8421053	0

```

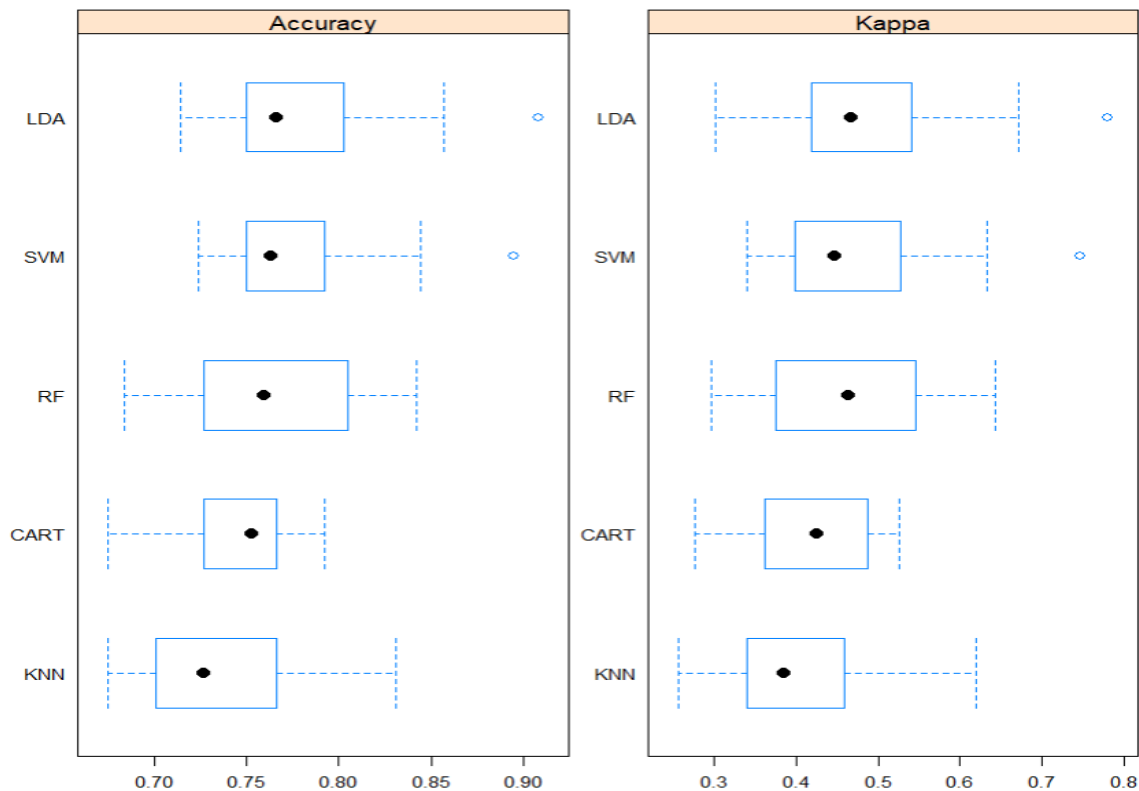
Kappa
```

	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.	NA's
CART	0.2762566	0.3620724	0.4241878	0.4151867	0.4861107	0.5250000	0
LDA	0.3011551	0.4192537	0.4662541	0.4862025	0.5308596	0.7812500	0
SVM	0.3391908	0.3997116	0.4460612	0.4621585	0.5234605	0.7475083	0
KNN	0.2553191	0.3406000	0.3841761	0.3984995	0.4539789	0.6195363	0
RF	0.2951613	0.3778304	0.4640696	0.4630809	0.5447483	0.6426332	0

CODE:


```
# box and whisker plots to compare models
```

```
scales <- list(x=list(relation="free"), y=list(relation="free"))
bwplot(results, scales=scales)
```

OUTPUT:**CODE:**

```
# density plots of accuracy
```

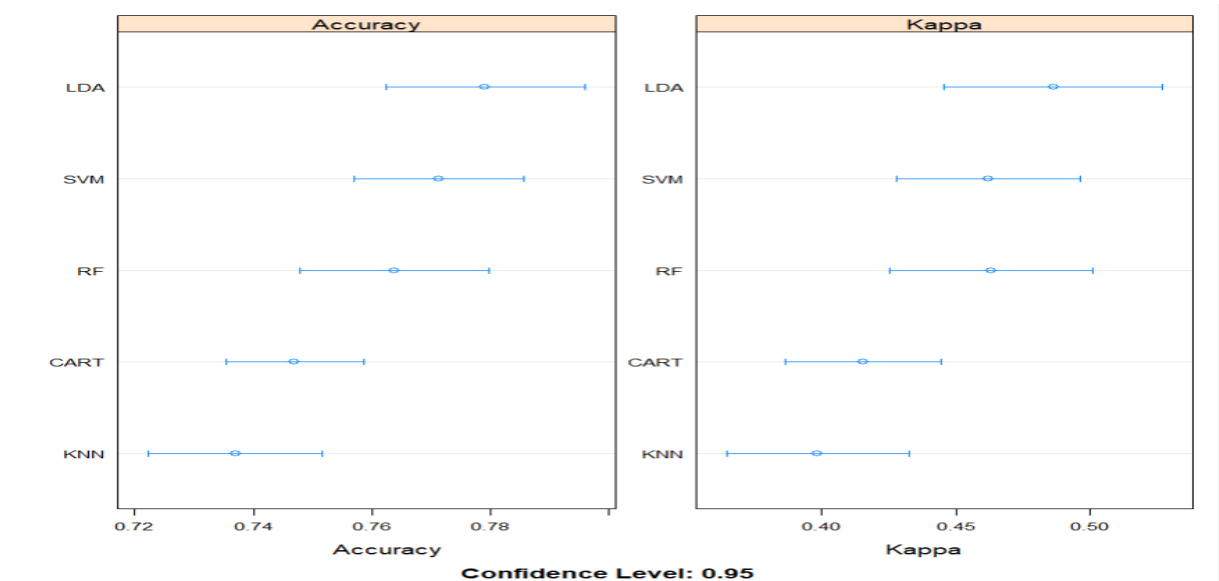
```
scales <- list(x=list(relation="free"), y=list(relation="free"))
densityplot(results, scales=scales, pch = "|")
```

OUTPUT:


CODE:**# dot plots of accuracy**

```
scales <- list(x=list(relation="free"), y=list(relation="free"))
```

```
dotplot(results, scales=scales)
```



OUTPUT:**CODE:**

```
parallelplot(results) # parallel plots to compare models
```

OUTPUT:**CODE:**

```
splom(results) # pair-wise scatterplots of predictions to compare models
```

```
xyplot(results, models=c("LDA", "SVM")) # xyplot plots to compare models
```

OUTPUT:**CODE:**

```
diffs <- diff(results) # difference in model predictions
```

```
summary(diffs) # summarize p-values for pair-wise comparisons
```

OUTPUT:

```
> summary(diffs)
```

```
Call:
```

```
summary.diff.resamples(object = diffs)
```

```
p-value adjustment: bonferroni
```

```
Upper diagonal: estimates of the difference
```

```
Lower diagonal: p-value for H0: difference = 0
```

```
Accuracy
```

	CART	LDA	SVM	KNN	RF
CART		-0.032137	-0.024322	0.010019	-0.016883
LDA	0.0011862		0.007815	0.042157	0.015254
SVM	0.0116401	0.9156892		0.034342	0.007439
KNN	1.0000000	6.68e-05	0.0002941		-0.026902
RF	0.2727542	0.4490617	1.0000000	0.0183793	

```
Kappa
```

	CART	LDA	SVM	KNN	RF
CART		-0.0710158	-0.0469717	0.0166872	-0.0478942
LDA	0.0008086		0.0240440	0.0877029	0.0231215
SVM	0.0258079	0.3562734		0.0636589	-0.0009225
KNN	1.0000000	0.0003858	0.0040823		-0.0645814
RF	0.0211763	1.0000000	1.0000000	0.0158974	

EXPERIMENT NO:12

AIM:Create aR program to implement Regression Comparison on Boston Housing dataset and compare and visualize the results

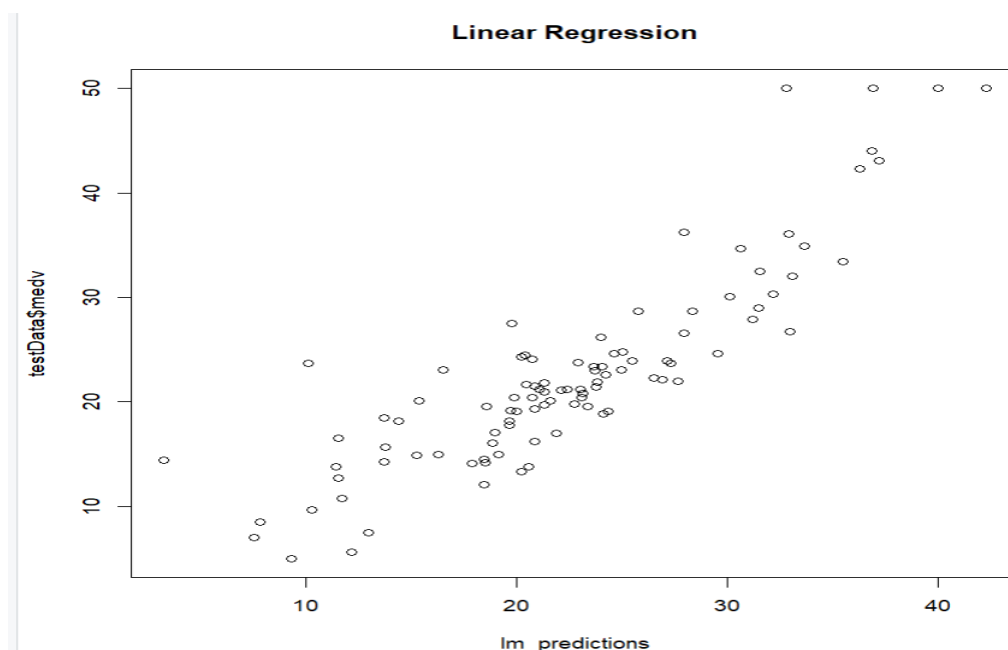
DESCRIPTION:The “**Boston**” Housing dataset is a commonly used dataset for regression tasks in machine learning. It contains various features related to housing in Boston, such as the crime rate, average number of rooms per dwelling, and more. The target variable in this dataset is typically the median value of owner-occupied homes (in thousands of dollars).

CODE:**# Load necessary libraries**

```
library(MASS)
library(caret)
library(e1071)
library(randomForest)
library(rpart)
library(rpart.plot)
data(Boston) # Load BostonHousing dataset
set.seed(123) # Set seed for reproducibility
# Split data into training and testing sets

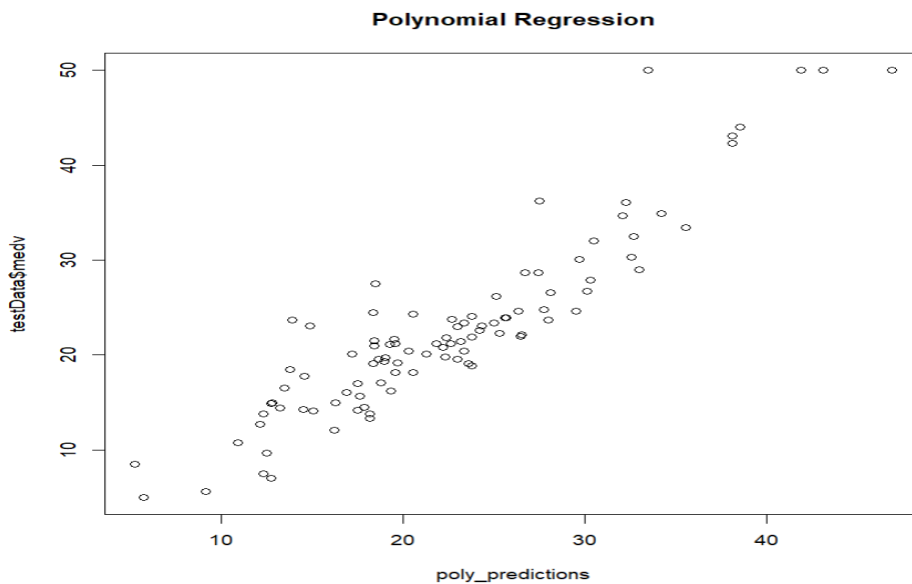
trainIndex<- createDataPartition(Boston$medv, p = 0.8, list = FALSE)
trainData<- Boston[trainIndex, ]
testData<- Boston[-trainIndex, ]
# Linear Regression

lm_model<- lm(medv ~ ., data = trainData)
lm_predictions<- predict(lm_model, newdata = testData)
plot(lm_predictions, testData$medv, main = "Linear Regression")
```

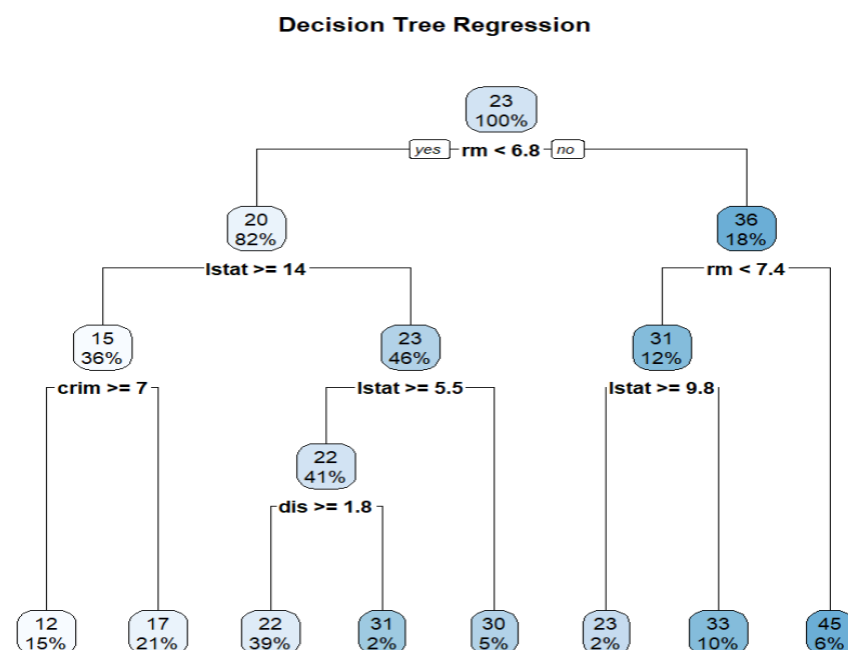
OUTPUT:

CODE:**# Polynomial Regression (excluding 'chas')**

```
poly_model<- lm(medv ~ poly(crim, 2) + poly(zn, 2) + poly(indus, 2) + poly(nox, 2) + poly(rm, 2) + poly(age, 2) +
poly(dis, 2) + poly(rad, 2) + poly(tax, 2) + poly(ptratio, 2) + poly(black, 2) + poly(lstat, 2), data = trainData)
poly_predictions<- predict(poly_model, newdata = testData)
plot(poly_predictions, testData$medv, main = "Polynomial Regression")
```

OUTPUT:**CODE:****# Decision Tree Regression**

```
tree_model<- rpart(medv ~ ., data = trainData, method = "anova")
tree_predictions<- predict(tree_model, newdata = testData)
rpart.plot(tree_model, main = "Decision Tree Regression")
```

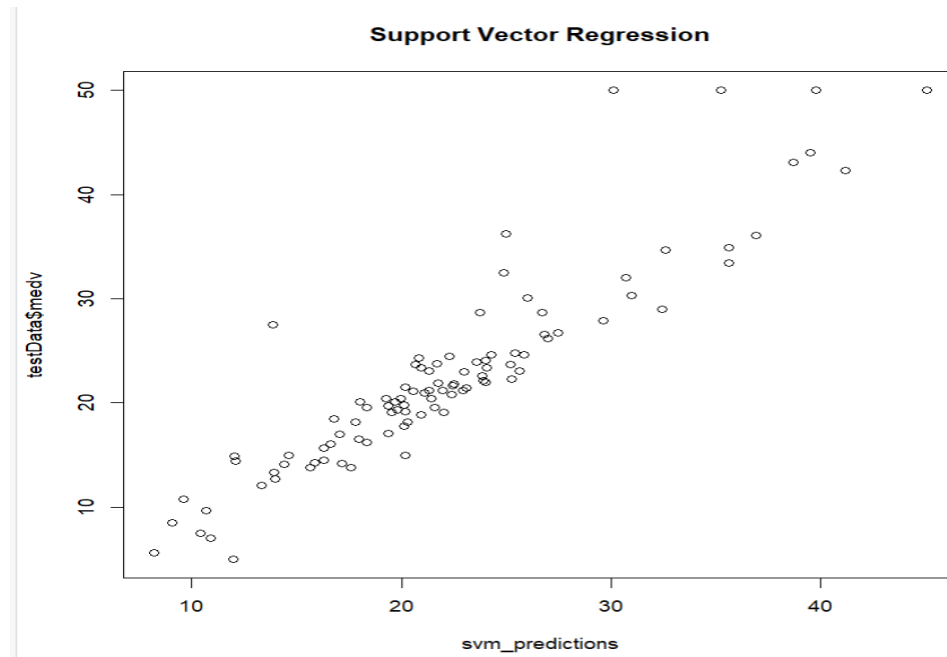
OUTPUT:

CODE:**# Support Vector Regression**

```
svm_model<- svm(medv ~ ., data = trainData)

svm_predictions<- predict(svm_model, newdata = testData)


plot(svm_predictions, testData$medv, main = "Support Vector Regression")
```

OUTPUT:**CODE:****# Random Forest Regression**

```
rf_model<- randomForest(medv ~ ., data = trainData)

rf_predictions<- predict(rf_model, newdata = testData)

plot(rf_predictions, testData$medv, main = "Random Forest Regression")
```

OUTPUT:

CODE:**# Model Summaries**

```
summary(lm_model)
```

```
summary(poly_model)
```

OUTPUT:

```
Call:
lm(formula = medv ~ ., data = trainData)

Residuals:
    Min       1Q   Median       3Q      Max
-14.9550  -2.7996  -0.4647   1.7767  25.0993

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)  37.733617   5.619935   6.714 6.63e-11 ***
crim         -0.093857   0.039157  -2.397 0.016999 *
zn           0.039436   0.015987   2.467 0.014062 *
indus        -0.012988   0.069595  -0.187 0.852059
chas          2.290187   0.940621   2.435 0.015346 *
nox          -17.130560   4.342272  -3.945 9.45e-05 ***
rm           3.499219   0.451445   7.751 7.87e-14 ***
age           0.009823   0.015510   0.633 0.526905
dis          -1.390769   0.230614  -6.031 3.77e-09 ***
rad           0.330939   0.077135   4.290 2.25e-05 ***
tax          -0.012386   0.004342  -2.852 0.004568 **
ptratio      -0.960676   0.150307  -6.391 4.66e-10 ***
black         0.009841   0.002935   3.353 0.000877 ***
lstat        -0.562095   0.059180  -9.498 < 2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 4.801 on 393 degrees of freedom
Multiple R-squared:  0.7346, Adjusted R-squared:  0.7258
F-statistic: 83.68 on 13 and 393 DF, p-value: < 2.2e-16
```

```
Call:
lm(formula = medv ~ poly(crim, 2) + poly(zn, 2) + poly(indus,
  2) + poly(nox, 2) + poly(rm, 2) + poly(age, 2) + poly(dis,
  2) + poly(rad, 2) + poly(tax, 2) + poly(ptratio, 2) + poly(black,
  2) + poly(lstat, 2), data = trainData)

Residuals:
    Min       1Q   Median       3Q      Max
-22.8962  -2.1719  -0.1371   1.8323  25.3764

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)  22.5106     0.1972 114.143 < 2e-16 ***
poly(crim, 2)1 -38.0606     7.3862  -5.153 4.12e-07 ***
poly(crim, 2)2  16.4345     6.4240   2.558 0.01090 *
poly(zn, 2)1  -3.3820     7.5876  -0.446 0.65605
poly(zn, 2)2   8.4794     4.3750   1.938 0.05334 .
poly(indus, 2)1  3.6202     8.5393   0.424 0.67185
poly(indus, 2)2  5.8043     6.3912   0.908 0.36436
poly(nox, 2)1 -56.7798    12.3064  -4.614 5.40e-06 ***
poly(nox, 2)2   3.2721     6.7800   0.483 0.62965
poly(rm, 2)1  43.9986     5.8369   7.538 3.50e-13 ***
poly(rm, 2)2  35.1658     4.8284   7.283 1.88e-12 ***
poly(age, 2)1   6.9322     8.1659   0.849 0.39645
poly(age, 2)2   3.5476     4.6824   0.758 0.44913
poly(dis, 2)1 -51.3516    10.0059  -5.132 4.57e-07 ***
poly(dis, 2)2  12.4765     5.7688   2.163 0.03118 *
poly(rad, 2)1  68.9626    22.0153   3.132 0.00187 **
poly(rad, 2)2  -5.2775     5.3401  -0.988 0.32365
poly(tax, 2)1 -40.7625    20.4724  -1.991 0.04718 *
poly(tax, 2)2   6.0232     8.3002   0.726 0.46848
poly(ptratio, 2)1 -35.9445     5.8536  -6.141 2.06e-09 ***
poly(ptratio, 2)2  11.9187     5.3280   2.237 0.02586 *
poly(black, 2)1  12.7193     4.7348   2.686 0.00754 **
poly(black, 2)2  -4.1461     4.2448  -0.977 0.32931
poly(lstat, 2)1 -88.7852     7.6135  -11.662 < 2e-16 ***
poly(lstat, 2)2  29.0043     5.2507   5.524 6.14e-08 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 3.979 on 382 degrees of freedom
Multiple R-squared:  0.8228, Adjusted R-squared:  0.8117
F-statistic: 73.91 on 24 and 382 DF, p-value: < 2.2e-16
```

CODE:

```
print(tree_model)
```

OUTPUT:

```
> print(tree_model)
n= 407

node), split, n, deviance, yval
* denotes terminal node

1) root 407 34125.5800 22.51057
 2) rm< 6.8375 334 12681.2200 19.62695
   4) lstat>=14.4 146 2607.9150 15.05822
     8) crim>=7.006285 60 659.0840 11.86000 *
     9) crim< 7.006285 86 906.9406 17.28953 *
   5) lstat< 14.4 188 4659.1330 23.17500
     10) lstat>=5.51 167 2966.0360 22.35329
        20) dis>=1.7548 160 1307.6640 21.98312 *
        21) dis< 1.7548 7 1135.3290 30.81429 *
     11) lstat< 5.51 21 683.6381 29.70952 *
   3) rm>=6.8375 73 5959.9890 35.70411
     6) rm< 7.443 49 2037.2070 31.28367
        12) lstat>=9.76 8 440.5750 23.02500 *
        13) lstat< 9.76 41 944.5190 32.89512 *
     7) rm>=7.443 24 1010.4700 44.72917 *
```


CODE:

```
summary(svm_model)

print(rf_model)
```

OUTPUT:

```
Call:
svm(formula = medv ~ ., data = trainData)

Parameters:
  SVM-Type:  eps-regression
  SVM-Kernel: radial
    cost:    1
  gamma:    0.07692308
  epsilon:  0.1

Number of Support Vectors: 269
```

```
Call:
randomForest(formula = medv ~ ., data = trainData)
  Type of random forest: regression
    Number of trees: 500
No. of variables tried at each split: 4

Mean of squared residuals: 11.59743
  % Var explained: 86.17
```

CODE:**# Calculate RMSE for each model**

```
lm_rmse<- rmse(testData$medv - lm_predictions)
poly_rmse<- rmse(testData$medv - poly_predictions)
tree_rmse<- rmse(testData$medv - tree_predictions)
svm_rmse<- rmse(testData$medv - svm_predictions)
rf_rmse<- rmse(testData$medv - rf_predictions)
```

Print RMSEs for each model

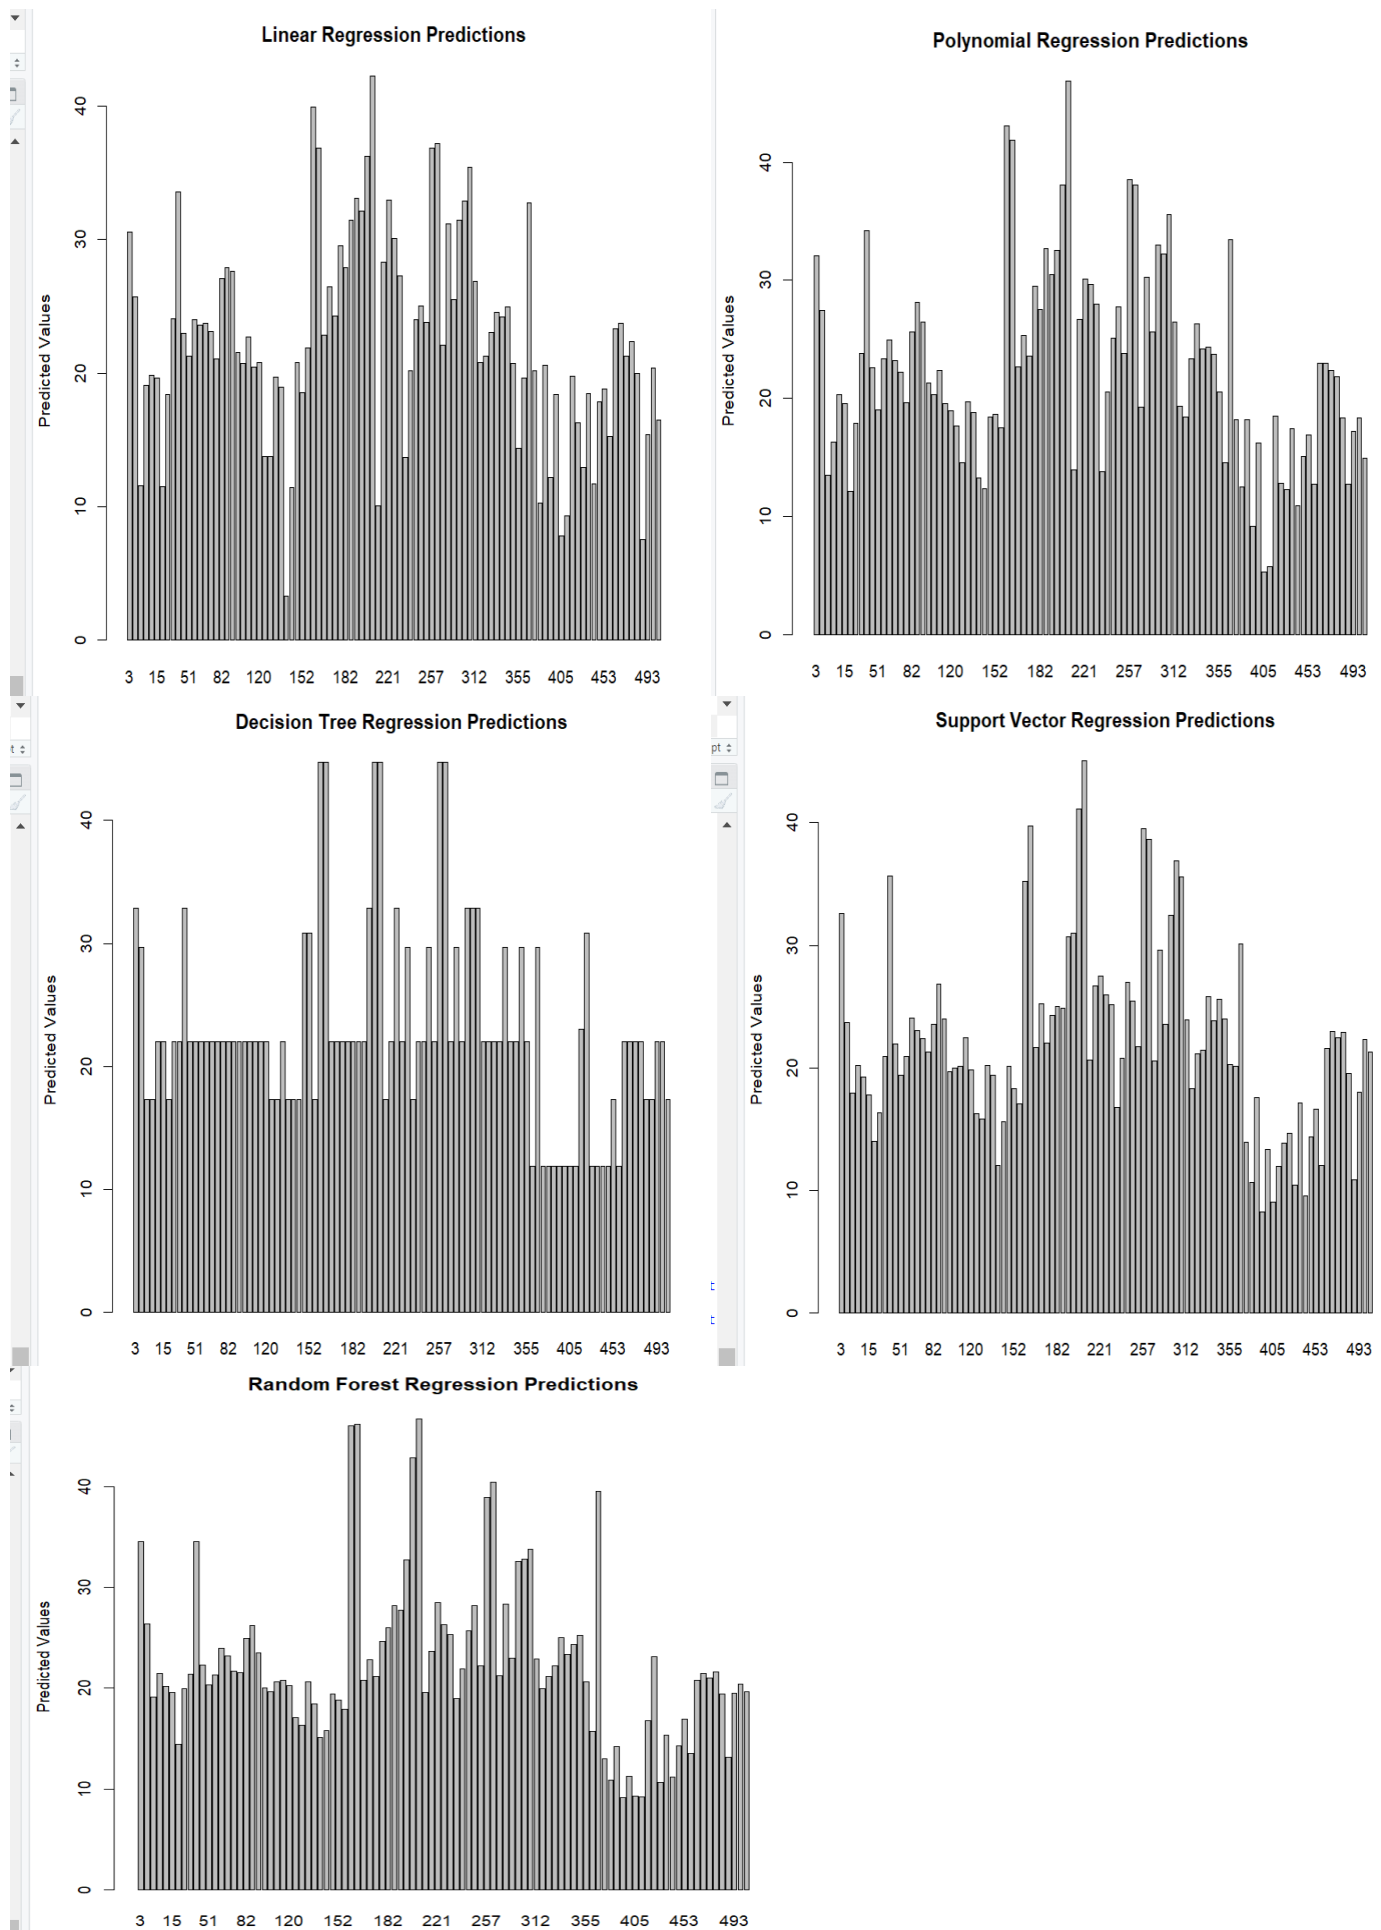
```
cat("Linear Regression RMSE:",lm_rmse,"\n")
cat("Polynomial Regression RMSE:",poly_rmse,"\n")
cat("Decision Tree Regression RMSE:",tree_rmse,"\n")
cat("Support Vector Regression RMSE:",svm_rmse,"\n")
cat("Random Forest Regression RMSE:",rf_rmse,"\n")
```

OUTPUT:

```
> cat("Linear Regression RMSE:", lm_rmse, "\n")
Linear Regression RMSE: 4.588948
> cat("Polynomial Regression RMSE:", poly_rmse, "\n")
Polynomial Regression RMSE: 3.770119
> cat("Decision Tree Regression RMSE:", tree_rmse, "\n")
Decision Tree Regression RMSE: 4.925522
> cat("Support Vector Regression RMSE:", svm_rmse, "\n")
Support Vector Regression RMSE: 3.91893
> cat("Random Forest Regression RMSE:", rf_rmse, "\n")
Random Forest Regression RMSE: 3.036031
```

CODE:**# Create bar graphs for each model's predictions**

```
models <- c("Linear Regression", "Polynomial Regression", "Decision Tree Regression", "Support Vector Regression",
"Random Forest Regression")
predictions <- list(lm_predictions, poly_predictions, tree_predictions, svm_predictions, rf_predictions)
for (i in 1:length(models)) {
  par(mfrow=c(1,1))
  barplot(predictions[[i]], main = paste(models[i], "Predictions"), ylab = "Predicted Values")}
```

OUTPUT:

EXPERIMENT NO:13

AIM:Create a R program to implement Clustering Comparison on College dataset and compare and visualize the results

DESCRIPTION:The "College" dataset is another commonly used dataset, typically for tasks like clustering and classification. This dataset contains information about various colleges in the United States. The dataset includes features related to the colleges, such as their names, types (public or private), number of students, acceptance rates, and more. One of the common uses of this dataset is for clustering colleges based on their characteristics.

CODE:**# Load required libraries**

```
library(ISLR2)
```

```
library(cluster)
```

```
library(dbSCAN)
```

```
library(factoextra)
```

```
library(mclust)
```

```
library(ggplot2)
```

```
library(kohonen)
```

```
library(e1071) # For FCM
```

```
library(reshape2) # For data manipulation
```

```
data("College") # Load the College dataset
```

Select relevant features for clustering (excluding the college name)

```
college_data<- College[, -1] # Exclude the first column (college name)
```

```
scaled_data<- scale(college_data)# Standardize the data
```

Create a function to calculate silhouette scores

```
calculate_silhouette<- function(method, data, clusters) {
```

```
  if (!is.null(clusters)) {
```

```
    silhouette_score<- mean(silhouette(clusters, dist(data)))
```

```
    return(silhouette_score)
```

```
  } else {
```

```
    return(NA)
```

```
  }}
```

Initialize a dataframe to store results

```
results <- data.frame(Method = character(0), Silhouette = numeric(0))
```

Perform K-Means clustering

```
kmeans_result<- kmeans(scaled_data, centers = 3, nstart = 25)
kmeans_silhouette<- calculate_silhouette("K-Means", scaled_data, kmeans_result$cluster)
results <- rbind(results, data.frame(Method = "K-Means", Silhouette = kmeans_silhouette))
```

Perform Hierarchical Clustering

```
hierarchical_result<- hclust(dist(scaled_data))
hierarchical_silhouette<- calculate_silhouette("Hierarchical", scaled_data, cutree(hierarchical_result, k = 3))
results <- rbind(results, data.frame(Method = "Hierarchical", Silhouette = hierarchical_silhouette))
```

Perform DBSCAN clustering

```
dbscan_result<- dbscan(scaled_data, eps = 3, MinPts = 5)
dbscan_silhouette<- calculate_silhouette("DBSCAN", scaled_data, dbscan_result$cluster)
results <- rbind(results, data.frame(Method = "DBSCAN", Silhouette = dbscan_silhouette))
```

Perform Gaussian Mixture Model (GMM) clustering

```
gmm_result<-Mclust(scaled_data,G=3)
```

OUTPUT:

```
> gmm_result<-Mclust(scaled_data,G=3)
fitting ...
=====| 100%
```

```
gmm_silhouette<- calculate_silhouette("GMM", scaled_data, gmm_result$classification)
results <- rbind(results, data.frame(Method = "GMM", Silhouette = gmm_silhouette))
```

Perform Self-Organizing Maps (SOM) clustering

```
som_result<- som(scale(scaled_data), grid = somgrid(3, 3, "hexagonal"))
som_silhouette<- calculate_silhouette("SOM", scale(scaled_data), som_result$unit.classif)
results <- rbind(results, data.frame(Method = "SOM", Silhouette = som_silhouette))
```

Perform Fuzzy C-Means (FCM) clustering

```
fcm_result<- cmeans(scaled_data, centers = 3, m = 1.2)
fcm_clusters<- fcm_result$cluster
fcm_silhouette<- calculate_silhouette("FCM", scaled_data, fcm_clusters)
results <- rbind(results, data.frame(Method = "FCM", Silhouette = fcm_silhouette))

Print(results)
```


OUTPUT:

```
> cat("SOM Cluster Summary:\n")
SOM Cluster Summary:
> print(som_result)
SOM of size 3x3 with a hexagonal topology.
Training data included.
> |
```

CODE:

```
cluster_assignments<- som_result$unit.classif# Get cluster assignments from the SOM result
```

```
unique_clusters<- unique(cluster_assignments)# Create a vector of unique cluster IDs
```



```
cluster_colors<- rainbow(length(unique_clusters))# Create a color palette for the clusters
```

```
assigned_colors<- cluster_colors[cluster_assignments]# Map cluster assignments to colors based on cluster IDs
```

```
# Visualize SOM Clustering with assigned colors
```

```
plot(som_result, type = "mapping", pchs = 20, col = assigned_colors, main = "SOM Clustering")
```

```
plot(scaled_data, col = fcm_clusters, pch = 19, main = "FCM Clustering") # Plot FCM clustering assignments
```

OUTPUT:**CODE:**

```
# Print the best method and its silhouette score
```

```
Cat("Best Clustering Method:",best_method,"\n")
```

```
Cat("Silhouette score:",results$Silhouette[best_method_index],"\n")
```

OUTPUT:

```
> plot(scaled_data, col = rgb(0.5, 0.5, 0.5), pch = 15, main = "K-Means Clustering")
> cat("Best Clustering Method:", best_method, "\n")
Best Clustering Method: SOM
> cat("Silhouette Score:", results$Silhouette[best_method_index], "\n")
Silhouette Score: 2.436695
```

CODE:

Create a bar plot of silhouette scores with different colors for each bar

```
ggplot(results, aes(x = Method, y = Silhouette, fill = Method)) +
```

```
geom_bar(stat = "identity") +
```

```
geom_text(aes(label = round(Silhouette, 2), vjust = -0.5), position = position_dodge(0.9)) + # Add labels on top of bars
```

```
labs(title = "Silhouette Scores for Clustering Algorithms", y = "Silhouette Score") +
```

```
theme_minimal() +
```

```
theme(axis.text.x = element_text(angle = 45, hjust = 1)) # Rotate x-axis labels for better readability
```

OUTPUT: